

Библиотеки Python для Data Science: продолжение

Итоговый проект на тему: "Credit Default"

Студент - Михайлов А.В., Kaggle User Name: avidclam, Display Name: Aleksandr Mikhailov

```
In [1]: import numpy as np
import pandas as pd

import matplotlib
from matplotlib import pyplot as plt
%matplotlib inline

import seaborn as sns

rc_params = {'figure.figsize': [11, 4.5], 'lines.linewidth': 2.5}

matplotlib.rcParams.update(rc_params)
sns.set_theme(style="whitegrid")
sns.set_context("notebook", font_scale=1.2, rc=rc_params)
```

```
In [2]: from scipy.stats import chi2_contingency, mannwhitneyu
```

```
In [3]: from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.metrics import accuracy_score, classification_report, plot_confusion_matrix, plot_precision_recall_
```

```
In [4]: import xgboost as xgb, lightgbm as lgbm, catboost as catb
```

Импорт данных

```
In [5]: TRAIN_DATASET_PATH = './data/credit-default/train.csv'
TEST_DATASET_PATH = './data/credit-default/test.csv'
```

```
In [6]: df = pd.read_csv(TRAIN_DATASET_PATH)
test_df = pd.read_csv(TEST_DATASET_PATH)
display(df.shape, test_df.shape)
```

```
(7500, 17)
(2500, 16)
```

Целевые классы

Credit Default - факт невыполнения кредитных обязательств (0 - погашен вовремя, 1 - просрочка)

```
In [7]: target = 'Credit Default'
FREQ_NAME = 'Chance of Default' # Обозначение частоты события "1". Ниже встретится сокращение CoD = Chance of D
target_counts = df[target].value_counts(normalize=True)
target_disbalance = target_counts.loc[0] / target_counts.loc[1]
print(target_counts, target_disbalance, sep='\n***\n')
```

```
0    0.718267
1    0.281733
Name: Credit Default, dtype: float64
****
2.5494557501183155
```

Вспомогательные функции

```
In [8]: # Возвращает частоту единичной целевой переменной для каждого значения категориального признака (пропуски считая)
def feature_target_counts(df, features, na_value='NA', target=target, freq_name=FREQ_NAME):
    result_df = df.filter(features + [target]).fillna(na_value).value_counts().unstack(fill_value=0)
    result_df['Count'] = result_df[0] + result_df[1]
    result_df['Share %'] = round(result_df['Count'] * 100 / df.shape[0], 2)
    result_df[freq_name] = result_df[1] / result_df['Count']
    return result_df
```

```
In [9]: # Возвращает частоту единичной целевой переменной для интервала распределения признака (на базе квантилей)
def get_quantile_freq(df, feature, nbins, target=target, freq_name=FREQ_NAME):
    _feature_df = df.filter([feature, target]).dropna()
    quantile_df = pd.DataFrame(_feature_df[feature].quantile(np.linspace(0, 1, nbins+1)))
    quantile_df[feature].iloc[[0, -1]] = np.NINF, np.inf
    quantile_df = quantile_df.reset_index().rename(columns = {'index': 'Quantile'})
    _feature_df['Feature Bin'] = pd.cut(_feature_df[feature], bins = quantile_df[feature], labels=False)
    return feature_target_counts(_feature_df, ['Feature Bin'], target=target, freq_name=freq_name).join(quantile
```

```
In [10]: # Используется для генерации target encoding-признака, сопоставляющего исходному значению признака
# частоту единичной целевой переменной соответствующего интервала распределения (из функции get_quantile_freq)
def get_feature_freq(df, quantile_freq, na_value=None, freq_name=FREQ_NAME):
    feature = quantile_freq.columns[-1]
    _feature_df = pd.DataFrame({'Feature Bin': pd.cut(df[feature], bins=quantile_freq[feature].tolist() + [np.in
result_df = pd.merge(_feature_df, quantile_freq.reset_index().filter(['Feature Bin', freq_name]), how='left'
if na_value is not None:
    result_df[freq_name].fillna(na_value, inplace=True)
return result_df[freq_name].values
```

Признаки: анализ существующих, генерация дополнительных

```
In [11]: feature_candidates = [] # Здесь будем собирать кандидатов в признаки модели
```

```
In [12]: # Предварительная значимость числовых признаков исходя из корреляции с целевой переменной
target_corr = df.corr()[target].abs().sort_values(ascending=False)[1:]
print(target_corr)
```

```
Credit Score          0.442050
Current Loan Amount   0.226522
Annual Income         0.101375
Number of Open Accounts 0.028884
Tax Liens             0.024368
Years of Credit History 0.020567
Number of Credit Problems 0.020088
Current Credit Balance 0.019522
Monthly Debt          0.018480
Maximum Open Credit   0.014275
Months since last delinquent 0.002598
Bankruptcies          0.001648
Name: Credit Default, dtype: float64
```

```
In [13]: # Список категориальных признаков:
set(df.columns) - set(target_corr.index) - {target}
```

```
Out[13]: {'Home Ownership', 'Purpose', 'Term', 'Years in current job'}
```

```
In [14]: # Признаки с пропущенными значениями
print(df.columns[df.apply(lambda x: x.isnull().any()).sort_values().tolist()])
```

```
['Annual Income', 'Bankruptcies', 'Credit Score', 'Months since last delinquent', 'Years in current job']
```

Credit Score

```
In [15]: feature_name = 'Credit Score'
df[feature_name].isnull().value_counts(normalize=True)
```

```
Out[15]: False    0.7924
         True     0.2076
         Name: Credit Score, dtype: float64
```

```
In [16]: # Сразу заметим, что пропущенные значения признака Credit Score совпадают с пропущенными значениями Annual Income
all(df['Credit Score'].isnull() == df['Annual Income'].isnull())
```

```
Out[16]: True
```

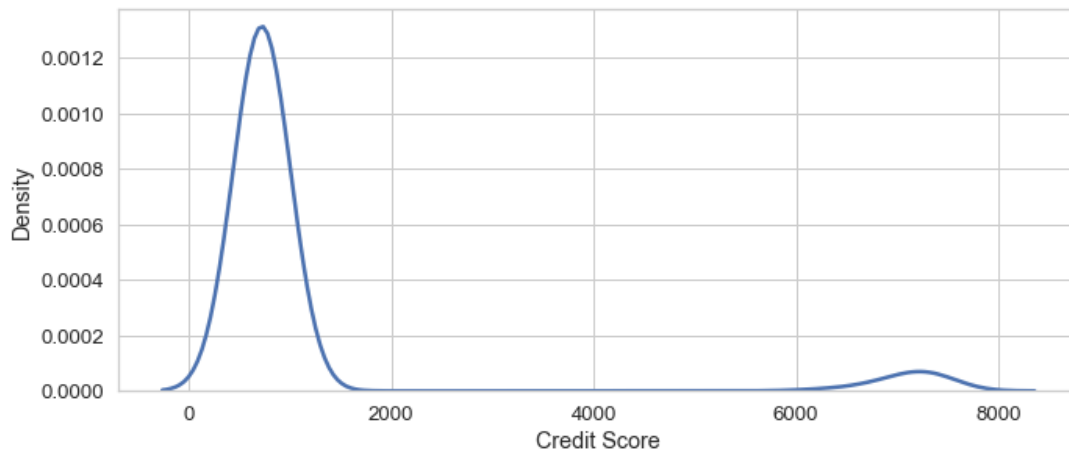
Примечание: вне данного ноутбука была сделана попытка заполнить пропущенные значения Credit Score и Annual Income через регрессию по другим признакам, но значения R2 в районе ~0.3 показали недостаточными для практического применения этого метода.

```
In [17]: # Запомним значение частоты дефолта для данных, у которых пропущен кредитный рейтинг и значения годового дохода,
credit_score_missing_freq = df.loc[df[feature_name].isnull(), target].mean()
```

```
credit_score_missing_freq
```

```
Out[17]: 0.3397559409120103
```

```
In [18]: sns.kdeplot(data=df, x=feature_name);
```




```
In [19]: # Выясним, как влияет "расположение на шкале" на целевую переменную
feature_df = df.assign(Tag=np.sign(df[feature_name] - 4000).map({-1: 'Low', 1: 'High'}))
feature_target_counts(feature_df, ['Tag', 'Missing']).sort_values(FREQ_NAME, ascending=False)
```

```
Out[19]: Credit Default    0    1  Count  Share %  Chance of Default
Tag
High      0    400    400    5.33    1.000000
Missing  1028    529   1557   20.76    0.339756
Low     4359   1184   5543   73.91    0.213603
```

Как видим, "десятикратные" показатели рейтинга однозначно свидетельствуют о ненадежности клиента.

```
In [20]: # Посмотрим, как частота дефолта меняется в квантилях распределения "нижней кривой" признака
qfreq_df = get_quantile_freq(df[df[feature_name] < 4000], feature_name, nbins=15)
sns.regplot(data=qfreq_df, x='Quantile', y=FREQ_NAME);
```



```
In [21]: # Создадим признак Credit Score CoD, в котором значению Credit Score сопоставлена частота единичной целевой пере
feature_freq_name = f"{feature_name} CoD"
df[feature_freq_name] = get_feature_freq(df, qfreq_df, na_value=credit_score_missing_freq)
test_df[feature_freq_name] = get_feature_freq(test_df, qfreq_df, na_value=credit_score_missing_freq)
```

```
In [22]: # Учет высокий риск "десятикратных" значений основного признака
df.loc[df[feature_name] > 4000, feature_freq_name] = 1.0
test_df.loc[test_df[feature_name] > 4000, feature_freq_name] = 1.0
```

```
In [23]: df.filter([feature_name, feature_freq_name]).corr()
```

```
Out[23]:
```

	Credit Score	Credit Score CoD
Credit Score	1.00000	0.91332
Credit Score CoD	0.91332	1.00000


```
In [24]: # Корреляция двух признаков большая, оставим только сгенерированный через целевую переменную
feature_candidates.append(feature_freq_name)
```

Annual Income - годовой доход

```
In [25]: feature_name = 'Annual Income'
df[feature_name].isnull().value_counts(normalize=True)
```

```
Out[25]: False    0.7924
         True     0.2076
         Name: Annual Income, dtype: float64
```

```
In [26]: # Создадим признак Annual Income CoD, аналогичный по построению Credit Score CoD
qfreq_df = get_quantile_freq(df, feature_name, nbins=9)
sns.regplot(data=qfreq_df, x='Quantile', y=FREQ_NAME)
feature_freq_name = f"{feature_name} CoD"
df[feature_freq_name] = get_feature_freq(df, qfreq_df, na_value=credit_score_missing_freq)
test_df[feature_freq_name] = get_feature_freq(test_df, qfreq_df, na_value=credit_score_missing_freq)
```



```
In [27]: df.filter([feature_name, feature_freq_name]).corr()
```

```
Out[27]:
```

	Annual Income	Annual Income CoD
Annual Income	1.000000	-0.826125
Annual Income CoD	-0.826125	1.000000

```
In [28]: feature_candidates.append(feature_freq_name)
```

Current Loan Amount - текущая сумма кредита

```
In [29]: feature_name = 'Current Loan Amount'
```

```
In [30]: # Есть странное значение 99999999, обозначающее какое-то особое состояние
(df[feature_name] == 99999999).value_counts()
```

```
Out[30]: False    6630
         True     870
         Name: Current Loan Amount, dtype: int64
```

```
In [31]: # Как можно убедиться, "особое состояние" = "уверенность в погашении кредита"
feature_df = df.assign(Tag=df[feature_name] == 99999999)
feature_target_counts(feature_df, ['Tag'])
```

```
Out[31]:
```

Credit Default	0	1	Count	Share %	Chance of Default
----------------	---	---	-------	---------	-------------------

Credit Default	0	1	Count	Share %	Chance of Default
False	4517	2113	6630	88.4	0.318703
True	870	0	870	11.6	0.000000

```
In [32]: # Создадим особый признак
new_feature_name = f"{feature_name} Low Risk"
df[new_feature_name] = (df[feature_name] == 9999999).astype(np.uint8)
test_df[new_feature_name] = (test_df[feature_name] == 9999999).astype(np.uint8)
```

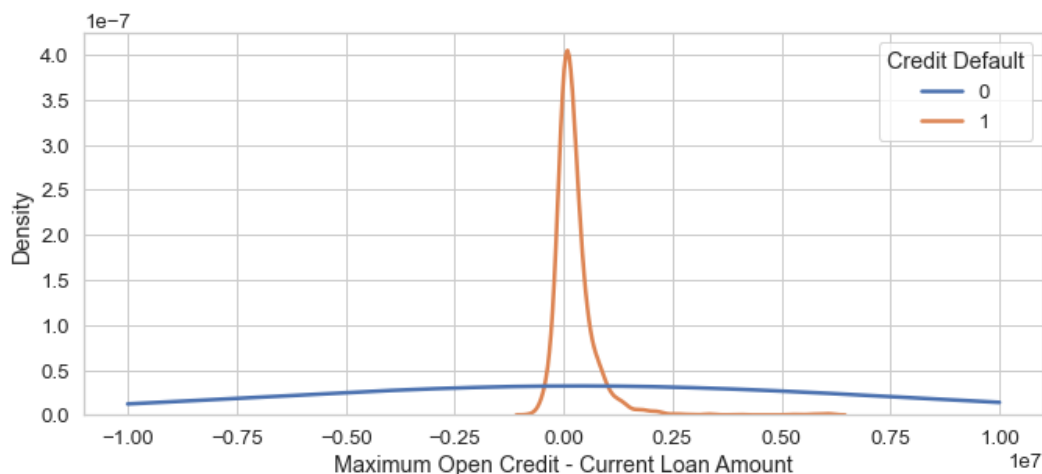
```
In [33]: feature_candidates.append(new_feature_name) # сам признак используем далее
```

Maximum Open Credit - наибольший открытый кредит

Сравнение с Current Loan Amount может быть полезно для разделения целевой переменной.

```
In [34]: feature_name = 'Maximum Open Credit'
cross_feature_name = 'Maximum Open Credit - Current Loan Amount'
df[cross_feature_name] = df[feature_name] - df['Current Loan Amount']
test_df[cross_feature_name] = test_df[feature_name] - test_df['Current Loan Amount']
```

```
In [35]: sns.kdeplot(data=df, x=cross_feature_name, hue=target, clip=(-1e7, 1e7));
```



```
In [36]: # Убедимся в отсутствии эквивалентности мат. ожиданий в исследуемых группах, с помощью критерия Манна-Уитни.
mannwhitneyu(
    df.loc[df[target] == 0, cross_feature_name].values,
    df.loc[df[target] == 1, cross_feature_name].values
)
```


```
Out[36]: MannwhitneyuResult(statistic=5371067.5, pvalue=7.317532151695133e-05)
```

```
In [37]: df.loc[(df['Current Loan Amount Low Risk'] == 1), cross_feature_name] = df.loc[(df['Current Loan Amount Low Risk'] == 1), cross_feature_name]
test_df.loc[(test_df['Current Loan Amount Low Risk'] == 1), cross_feature_name] = test_df.loc[(test_df['Current Loan Amount Low Risk'] == 1), cross_feature_name]
```

```
In [38]: feature_candidates.append(cross_feature_name)
```

Current Credit Balance - текущий кредитный баланс и Monthly Debt - ежемесячный долг

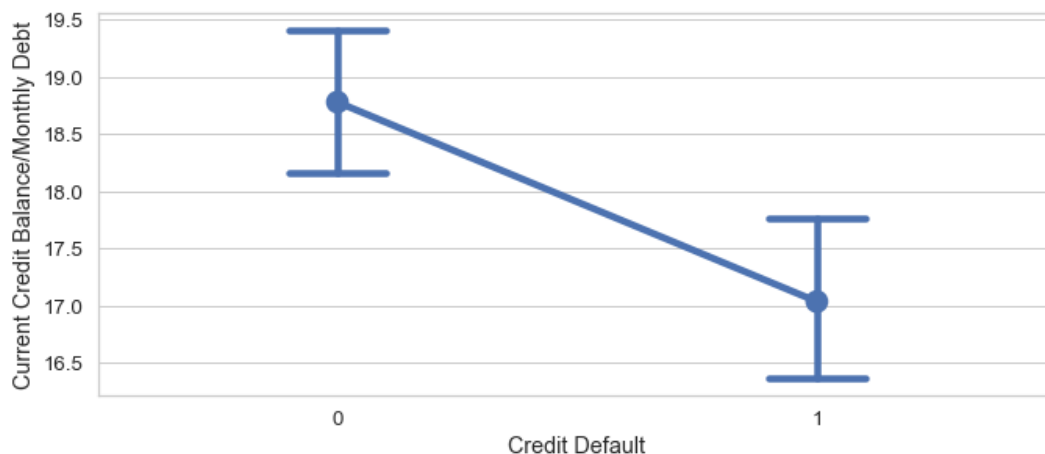
```
In [39]: # Оба признака не особенно эффективно разделяют целевую переменную
fig, ax = plt.subplots(1,2)
sns.pointplot(data=df, x=target, y='Current Credit Balance', capsize=.2, ax=ax[0])
sns.pointplot(data=df, x=target, y='Monthly Debt', capsize=.2, ax=ax[1])
fig.tight_layout()
```



```
In [40]: # Создадим синтетический признак частного от деления 'Current Credit Balance' на 'Monthly Debt'
df.loc[df['Monthly Debt'] == 0, 'Monthly Debt'] = np.NaN
df['Current Credit Balance/Monthly Debt'] = df['Current Credit Balance'] / df['Monthly Debt']
df['Current Credit Balance/Monthly Debt'].fillna(df['Current Credit Balance/Monthly Debt'].max(), inplace=True)
df['Monthly Debt'].fillna(0, inplace=True)
```

```
In [41]: test_df.loc[df['Monthly Debt'] == 0, 'Monthly Debt'] = np.NaN
test_df['Current Credit Balance/Monthly Debt'] = test_df['Current Credit Balance'] / test_df['Monthly Debt']
test_df['Current Credit Balance/Monthly Debt'].fillna(test_df['Current Credit Balance/Monthly Debt'].max(), inplace=True)
test_df['Monthly Debt'].fillna(0, inplace=True)
```

```
In [42]: # Синтетический признак разделяет немного лучше
sns.pointplot(data=df, x=target, y='Current Credit Balance/Monthly Debt', capsize=.2);
```



```
In [43]: feature_candidates.append('Current Credit Balance/Monthly Debt')
```

Tax Liens - налоговые обременения

```
In [44]: feature_name = 'Tax Liens'
agg_df = feature_target_counts(df, [feature_name])
agg_df
```

```
Out[44]: Credit Default    0    1  Count  Share %  Chance of Default
```

Tax Liens					
0.0	5301	2065	7366	98.21	0.280342
1.0	59	24	83	1.11	0.289157
2.0	15	15	30	0.40	0.500000
3.0	5	5	10	0.13	0.500000
4.0	3	3	6	0.08	0.500000
5.0	1	1	2	0.03	0.500000

Credit Default	0	1	Count	Share %	Chance of Default
Tax Liens					
6.0	2	0	2	0.03	0.000000
7.0	1	0	1	0.01	0.000000

```
In [45]: # Группы мало пригодны для разделения целевой переменной
chi2, p, dof, ex = chi2_contingency(agg_df.loc[[0,1], [0,1]])
p
```

Out[45]: 0.956212700856196

Number of Credit Problems - количество проблем с кредитом

```
In [46]: feature_name = 'Number of Credit Problems'
agg_df = feature_target_counts(df, [feature_name])
agg_df
```

```
Out[46]:
```

	Credit Default	0	1	Count	Share %	Chance of Default
Number of Credit Problems						
0.0	4654	1815	6469	86.25	0.280569	
1.0	641	241	882	11.76	0.273243	
2.0	60	33	93	1.24	0.354839	
3.0	21	14	35	0.47	0.400000	
4.0	3	6	9	0.12	0.666667	
5.0	3	4	7	0.09	0.571429	
6.0	4	0	4	0.05	0.000000	
7.0	1	0	1	0.01	0.000000	

```
In [47]: chi2, p, dof, ex = chi2_contingency(agg_df.loc[[0,1], [0,1]])
p
```

Out[47]: 0.6783214582441461

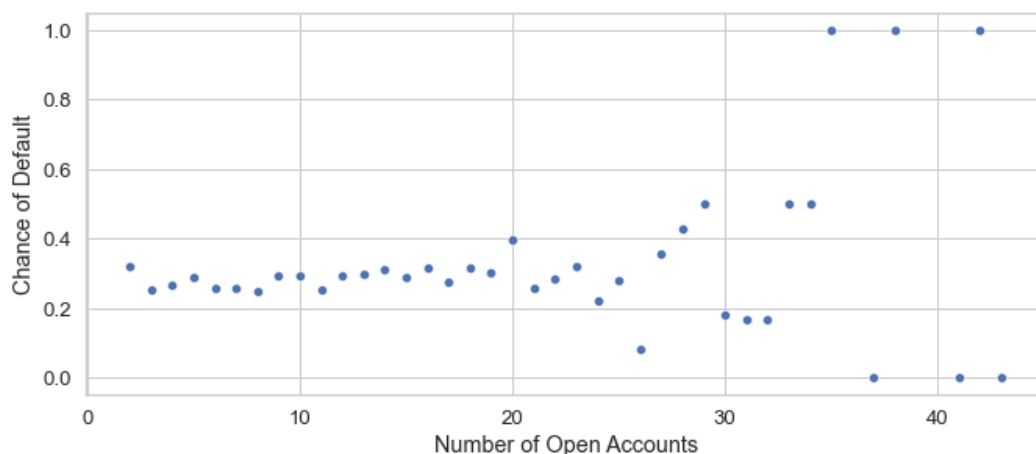
```
In [48]: chi2, p, dof, ex = chi2_contingency(agg_df.loc[[0,2], [0,1]])
p
```

Out[48]: 0.14293426550821686

Такая же история, группы либо малочисленные, либо не разделяют целевую переменную в достаточной мере.

Number of Open Accounts - количество открытых счетов

```
In [49]: feature_name = 'Number of Open Accounts'
sns.scatterplot(data=feature_target_counts(df, [feature_name]), x=feature_name, y=FREQ_NAME);
```



```
In [50]: # Начиная с некоторого значения признака, в данном случае, с 13, заметно отличие частоты дефолта
agg_df = feature_target_counts(df.assign(Over_12 =(df[feature_name] > 12)), ['Over_12'])
agg_df
```

```
Out[50]: Credit Default    0    1  Count  Share %  Chance of Default
Over_12
```

False	3692	1382	5074	67.65	0.272369
True	1695	731	2426	32.35	0.301319

```
In [51]: # Убедимся, что частоты не согласованы
chi2, p, dof, ex = chi2_contingency(agg_df.iloc[[0, 1], [0,1]])
p
```

```
Out[51]: 0.009886126527020072
```

```
In [52]: new_feature_name = f"{feature_name} Over 12"
df[new_feature_name] = (df[feature_name] > 12).astype(np.uint8)
test_df[new_feature_name] = (test_df[feature_name] > 12).astype(np.uint8)
```

```
In [53]: feature_candidates.append(new_feature_name)
```

Years in current job - количество лет на текущем месте работы

```
In [54]: feature_name = 'Years in current job'
feature_target_counts(df, [feature_name]).sort_values(FREQ_NAME, ascending=False)
```

```
Out[54]: Credit Default    0    1  Count  Share %  Chance of Default
```

Years in current job					
NA	234	137	371	4.95	0.369272
6 years	302	124	426	5.68	0.291080
< 1 year	404	159	563	7.51	0.282416
1 year	362	142	504	6.72	0.281746
5 years	371	145	516	6.88	0.281008
7 years	285	111	396	5.28	0.280303
4 years	338	131	469	6.25	0.279318
10+ years	1688	644	2332	31.09	0.276158
3 years	450	170	620	8.27	0.274194
2 years	512	193	705	9.40	0.273759
8 years	247	92	339	4.52	0.271386
9 years	194	65	259	3.45	0.250965


```
In [55]: # Создадим признак пропущенного значения, а значения самого признака не будем использовать
feature_missing_name = f"{feature_name} Missing"
df[feature_missing_name] = df[feature_name].isnull().astype(np.uint8)
test_df[feature_missing_name] = test_df[feature_name].isnull().astype(np.uint8)
```

```
In [56]: feature_candidates.append(feature_missing_name)
```

Years of Credit History - количество лет кредитной истории

```
In [57]: feature_name = 'Years of Credit History'
```

```
In [58]: sns.histplot(data=df, x=feature_name, hue=target);
```

```
In [59]: mannwhitneyu(
          df.loc[df[target] == 0, feature_name].values,
          df.loc[df[target] == 1, feature_name].values
        )
```

```
Out[59]: MannwhitneyuResult(statistic=5461129.5, pvalue=0.003171215731103466)
```

```
In [60]: feature_candidates.append(feature_name)
```

Months since last delinquent - количество месяцев с последней просрочки платежа

```
In [61]: feature_name = 'Months since last delinquent'
          df[feature_name].isnull().value_counts(normalize=True)
```

```
Out[61]: True      0.544133
          False    0.455867
          Name: Months since last delinquent, dtype: float64
```

```
In [62]: feature_df = df.filter([feature_name, 'Annual Income', target]).copy()
          feature_df['Tag'] = 10 * feature_df['Annual Income'].isnull() + feature_df[feature_name].isnull()
          feature_target_counts(feature_df, ['Tag'])
```

```
Out[62]: Credit Default    0    1  Count  Share %  Chance of Default
          Tag
          0  2003  733   2736    36.48    0.267909
          1  2356  851   3207    42.76    0.265357
          10  433  250    683     9.11    0.366032
          11  595  279    874    11.65    0.319222
```

```
In [63]: # Отметим несколько большую частоту дефолта в данных, у которых отсутствует значение Annual Income, но ненулевое
          new_feature_name = f"{feature_name} Special"
          df[new_feature_name] = ((10 * df['Annual Income'].isnull() + df[feature_name].isnull() == 10)).astype(np.uint8)
          test_df[new_feature_name] = (10 * df['Annual Income'].isnull() + test_df[feature_name].isnull() == 10).astype(np
```

```
In [64]: feature_target_counts(df, [new_feature_name])
```

```
Out[64]: Credit Default    0    1  Count  Share %  Chance of Default
          Months since last delinquent Special
          0  4954  1863   6817    90.89    0.273287
          1   433   250    683     9.11    0.366032
```

```
In [65]: feature_candidates.append(new_feature_name)
```

Bankruptcies - банкротства

```
In [66]: # Группы не особо помогают разделить целевую переменную
feature_target_counts(df, ['Bankruptcies'], -1)
```

```
Out[66]: Credit Default    0    1  Count  Share %  Chance of Default
Bankruptcies
-1.0    10    4    14    0.19    0.285714
 0.0  4782  1878   6660   88.80    0.281982
 1.0    569   217    786   10.48    0.276081
 2.0    20    11    31    0.41    0.354839
 3.0     5     2     7    0.09    0.285714
 4.0     1     1     2    0.03    0.500000
```

Purpose / Term / Home Ownership

```
In [67]: # Создадим новый признак, принимающий значения частоты дефолта, в зависимости от сочетания значений Purpose, Ter
new_feature_name = 'Purpose/Term/Home Ownership CoD'
```

```
In [68]: agg_df = (
    feature_target_counts(df, ['Purpose', 'Term', 'Home Ownership'])
    .sort_values('Share %', ascending=False)
    .filter(['Share %', FREQ_NAME])
    .reset_index()
    .rename(columns={FREQ_NAME: new_feature_name, 'Share %': f"{new_feature_name} Share"})
)
agg_df.head(7)
```

```
Out[68]: Credit Default    Purpose    Term    Home Ownership    Purpose/Term/Home Ownership CoD Share    Purpose/Term/Home Ownership CoD
0      debt consolidation    Short Term    Rent    27.35    0.268649
1      debt consolidation    Short Term    Home Mortgage    26.35    0.188259
2      debt consolidation    Long Term    Home Mortgage    11.84    0.382883
3      debt consolidation    Long Term    Rent    7.16    0.472998
4      debt consolidation    Short Term    Own Home    4.80    0.230556
5      other    Short Term    Rent    3.77    0.265018
6      home improvements    Short Term    Home Mortgage    2.96    0.171171
```

```
In [69]: df = pd.merge(df, agg_df, how='left')
df[new_feature_name].fillna(credit_score_missing_freq, inplace=True)
test_df = pd.merge(test_df, agg_df, how='left')
test_df[new_feature_name].fillna(credit_score_missing_freq, inplace=True)
```

```
In [70]: feature_candidates.append(new_feature_name)
```

Базовая модель

Корреляция признаков-кандидатов с целевой переменной и между собой

```
In [71]: df.filter(feature_candidates + [target]).corr()[target].abs().sort_values(ascending=False)[1:]
```

```
Out[71]: Credit Score CoD    0.428685
Purpose/Term/Home Ownership CoD    0.232629
Current Loan Amount Low Risk    0.226871
Annual Income CoD    0.119038
Months since last delinquent Special    0.059316
Years in current job Missing    0.044393
```

```
Current Credit Balance/Monthly Debt      0.036132
Number of Open Accounts Over 12         0.030106
Years of Credit History                  0.020567
Maximum Open Credit - Current Loan Amount 0.016367
Name: Credit Default, dtype: float64
```

```
In [72]: feature_corr_matrix = df.filter(feature_candidates).corr().round(2)
```

```
In [73]: sns_plot = sns.heatmap(feature_corr_matrix, annot=True, cmap='GnBu', annot_kws={"fontsize":10})
sns_plot.figure.set_size_inches((12, 6))
```



```
In [74]: # Считаем, что все кандидаты годятся
use_features = feature_candidates
```

Наборы данных для моделей

```
In [75]: X = df.filter(use_features)
y = df[target]
```

```
In [76]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, shuffle=True, stratify=y, random_state=X_train.shape, X_test.shape)
```

```
Out[76]: ((5250, 10), (2250, 10))
```

```
In [77]: basic_model = DecisionTreeClassifier(random_state=2020, max_depth=5)
```

Балансировка

Сравним методы Tomeklinks, SMOTE и задание весов классов базовой модели

Tomeklinks

```
In [78]: from imblearn.under_sampling import TomekLinks

t1 = TomekLinks()
X_train_balanced, y_train_balanced = t1.fit_sample(X_train, y_train)
```

```
y_train_balanced.value_counts()
```

```
Out[78]: 0    3100
         1    1479
         Name: Credit Default, dtype: int64
```

```
In [79]: basic_model.fit(X_train_balanced, y_train_balanced)
         print(classification_report(y_test, basic_model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.77	0.98	0.86	1616
1	0.81	0.25	0.39	634
accuracy			0.77	2250
macro avg	0.79	0.61	0.62	2250
weighted avg	0.78	0.77	0.73	2250

SMOTE

```
In [80]: from imblearn.over_sampling import SMOTE

         smote = SMOTE()
         X_train_balanced, y_train_balanced = smote.fit_sample(X_train, y_train)

         y_train_balanced.value_counts()
```

```
Out[80]: 1    3771
         0    3771
         Name: Credit Default, dtype: int64
```

```
In [81]: basic_model.fit(X_train_balanced, y_train_balanced)
         print(classification_report(y_test, basic_model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.82	0.73	0.77	1616
1	0.46	0.59	0.52	634
accuracy			0.69	2250
macro avg	0.64	0.66	0.65	2250
weighted avg	0.72	0.69	0.70	2250

Задание весов классов

```
In [82]: basic_model = DecisionTreeClassifier(random_state=2020, max_depth=5, class_weight={0:1, 1:target_disbalance})
         basic_model.fit(X_train, y_train)
         pred_train = basic_model.predict(X_train)
         pred_test = basic_model.predict(X_test)
         print(f'Accuracy на трейне {accuracy_score(y_train, pred_train)}')
         print(f'Accuracy на тесте {accuracy_score(y_test, pred_test)}')
         print(classification_report(y_train, basic_model.predict(X_train)))
         print(classification_report(y_test, basic_model.predict(X_test)))
```

```
Accuracy на трейне 0.7346666666666667
Accuracy на тесте 0.7368888888888889
```

	precision	recall	f1-score	support
0	0.82	0.81	0.81	3771
1	0.53	0.54	0.53	1479
accuracy			0.73	5250
macro avg	0.67	0.68	0.67	5250
weighted avg	0.74	0.73	0.74	5250

	precision	recall	f1-score	support
0	0.82	0.82	0.82	1616
1	0.53	0.53	0.53	634
accuracy			0.74	2250
macro avg	0.67	0.67	0.67	2250
weighted avg	0.74	0.74	0.74	2250

Далее будем использовать веса классов

```
In [83]: basic_model.fit(X, y)
basic_df = pd.DataFrame({'Id': test_df.index.values, target: basic_model.predict(test_df.filter(use_features))})
basic_df.to_csv('avidclam_submit_basic.csv', index=False)
```

Kaggle Score = 0.48945

Модели бустинга

Сразу зададим "улучшенные" параметры

XGBoost

```
In [84]: %%time
model_xgb = xgb.XGBClassifier(
    booster='gbtree',
    tree_method='hist',
    importance_type='total_gain',
    random_state=2020,
    scale_pos_weight=target_disbalance,
    learning_rate=0.07,
    max_depth=5,
    n_estimators=100,
    reg_lambda=30
)
model_xgb.fit(X_train, y_train)
print(classification_report(y_train, model_xgb.predict(X_train)))
print(classification_report(y_test, model_xgb.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.87	0.75	0.81	3771
1	0.53	0.71	0.60	1479
accuracy			0.74	5250
macro avg	0.70	0.73	0.70	5250
weighted avg	0.77	0.74	0.75	5250

	precision	recall	f1-score	support
0	0.84	0.73	0.78	1616
1	0.49	0.64	0.55	634
accuracy			0.71	2250
macro avg	0.66	0.69	0.67	2250
weighted avg	0.74	0.71	0.72	2250

Wall time: 229 ms

LightGBM

```
In [85]: %%time
model_lgbm = lgbm.LGBMClassifier(
    random_state=2020,
    class_weight={0:1, 1:target_disbalance},
    learning_rate=0.07,
    max_depth=5,
    n_estimators=100,
    reg_lambda=30,
    num_leaves=30
)
model_lgbm.fit(X_train, y_train)
print(classification_report(y_train, model_lgbm.predict(X_train)))
print(classification_report(y_test, model_lgbm.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.86	0.75	0.80	3771
1	0.52	0.70	0.60	1479
accuracy			0.73	5250
macro avg	0.69	0.73	0.70	5250
weighted avg	0.77	0.73	0.74	5250

	precision	recall	f1-score	support
0	0.84	0.74	0.78	1616
1	0.49	0.64	0.55	634
accuracy			0.71	2250
macro avg	0.66	0.69	0.67	2250
weighted avg	0.74	0.71	0.72	2250

Wall time: 201 ms

CatBoost

In [86]:

```
%%time
frozen_params = {
    'silent': True,
    'random_state': 2020,
    'eval_metric': 'F1',
    'l2_leaf_reg': 50,
    'class_weights': [1, target_disbalance],
    'iterations': 100,
    'max_depth': 5
}
model_cb = catb.CatBoostClassifier(**frozen_params)
model_cb.fit(X_train, y_train)
print(classification_report(y_train, model_cb.predict(X_train)))
print(classification_report(y_test, model_cb.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.85	0.72	0.78	3771
1	0.48	0.66	0.56	1479
accuracy			0.71	5250
macro avg	0.66	0.69	0.67	5250
weighted avg	0.74	0.71	0.72	5250

	precision	recall	f1-score	support
0	0.83	0.72	0.77	1616
1	0.47	0.63	0.54	634
accuracy			0.69	2250
macro avg	0.65	0.68	0.65	2250
weighted avg	0.73	0.69	0.71	2250

Wall time: 470 ms

Подбор гиперпараметров

Ни одна из моделей радикально не опередила другие по метрикам, для подбора гиперпараметров возьмем LightGBM, так как она имеет преимущество в скорости.

```
model = lgbm.LGBMClassifier(random_state=2020) param_grid={ 'class_weight': [{0:1, 1:target_disbalance*k} for k in (0.9, 1)], 'learning_rate': [0.07, 0.1], 'max_depth': [5, 6], 'n_estimators': [70, 100, 200], 'reg_lambda': [20, 30, 50], 'num_leaves': [20, 30, 70], } cv = StratifiedKFold(n_splits=3, random_state=31, shuffle=True) grid_search = GridSearchCV(model, param_grid, scoring="f1", cv=cv, n_jobs=-1, verbose=1) grid_search.fit(X_train, y_train) best_model = grid_search.best_estimator_ print(classification_report(y_test, best_model.predict(X_test)))
```

Лучшие параметры: {'class_weight': {0: 1, 1: 2.294510175106484}, 'learning_rate': 0.07, 'max_depth': 5, 'n_estimators': 200, 'num_leaves': 70, 'reg_lambda': 20} Метрики на трейне: precision recall f1-score support 0 0.87 0.81 0.84 3771 1 0.60 0.69 0.64 1479 accuracy 0.78 5250 macro avg 0.73 0.75 0.74 5250 weighted avg 0.79 0.78 0.79 5250

Метрики на тесте:

precision recall f1-score support 0 0.82 0.78 0.80 1616 1 0.51 0.58 0.54 634 accuracy 0.72 2250 macro avg 0.67 0.68 0.67 2250 weighted avg 0.74 0.72 0.73 2250

Финальная модель и экспорт данных

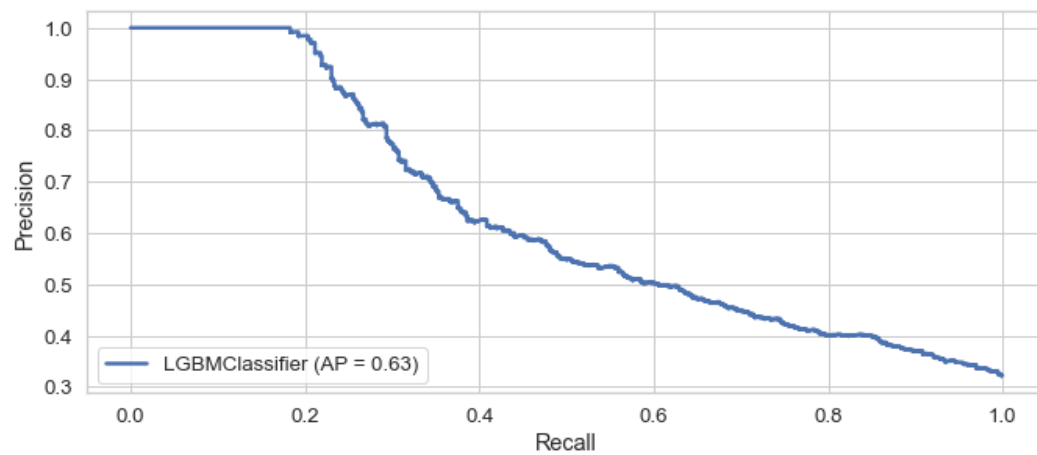
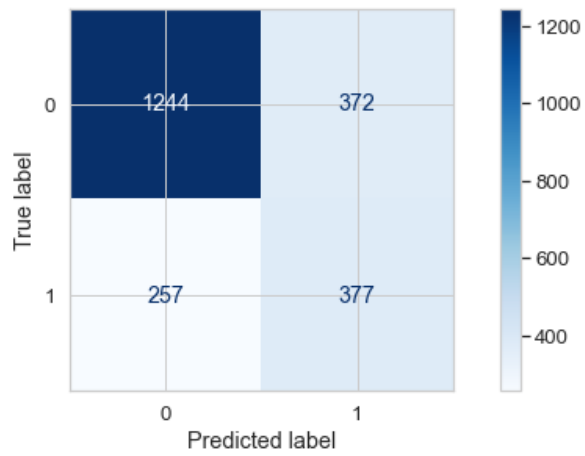
В финальной модели немного скорректированы параметры для снижения переобучения и повышения Precision (балансом Precision/Recall можно управлять, меняя class_weight).

In [87]:

```
final_params = {'class_weight': {0: 1, 1: target_disbalance*.92},
                'learning_rate': 0.07,
                'max_depth': 5,
                'n_estimators': 200,
                'num_leaves': 20,
                'reg_lambda': 50}
final_model = lgbm.LGBMClassifier(random_state=2020, **final_params)
final_model.fit(X_train, y_train)
print(classification_report(y_train, final_model.predict(X_train)))
print(classification_report(y_test, final_model.predict(X_test)))
plot_confusion_matrix(final_model, X_test, y_test, cmap=plt.cm.Blues)
plot_precision_recall_curve(final_model, X_test, y_test);
```

	precision	recall	f1-score	support
0	0.86	0.79	0.83	3771
1	0.56	0.68	0.62	1479
accuracy			0.76	5250

macro avg	0.71	0.74	0.72	5250
weighted avg	0.78	0.76	0.77	5250
	precision	recall	f1-score	support
0	0.83	0.77	0.80	1616
1	0.50	0.59	0.55	634
accuracy			0.72	2250
macro avg	0.67	0.68	0.67	2250
weighted avg	0.74	0.72	0.73	2250



```
In [88]: final_model.fit(X, y)
final_df = pd.DataFrame({'Id': test_df.index.values, target: final_model.predict(test_df.filter(use_features))})
final_df.to_csv('avidclam_submit.csv', index=False)
```

Kaggle Score: 0.54244